

The Value of Flexibility: Teleworkability, Sorting, and the Work-From-Home Wage Gap

Mitchell Valdes-Bobes

Anna Lukianova

April 18, 2025

Remote workers consistently earn higher wages than their on-site counterparts, a premium that persists even within detailed occupations and has evolved over time. We document these stylized facts using ACS, O*NET, and BLS data, constructing a novel occupation-level teleworkability index via machine learning. To understand these patterns, we develop a directed search model where workers (heterogeneous in skill) search for jobs which they value by the offered bundle of remote flexibility and wage, and firms (heterogeneous in remote work efficiency) optimally choose this bundles to post vacancies different submarkets. We estimate production function parameters and key distributions using industry-level data and calibrate remaining parameters. The model generates endogenous sorting, matching high-skill workers with high-remote-efficiency firms offering greater utility, including higher wages and more remote work. Our framework provides a structural explanation for the WFH premium driven by sorting and heterogeneity in firm remote work capabilities.

1 Introduction

In the United States, working from home is associated with higher real hourly wages (Gariety and Shaffer 2007), (Holgersen, Jia, and Svenkerud 2021). From 2013 to 2019, the real wage ratio of WFH employees to non-WFH employees remained relatively stable, fluctuating slightly between 1.35 and 1.40. This suggests that on average, remote employees earned about 35–40% more than their on-site counterparts during these years. However, in 2020, coinciding with the onset of the COVID-19 pandemic and the associated shift to remote work due to health restrictions, the wage ratio began to increase significantly. The ratio peaked in 2021 at approximately 1.6, indicating that the average real hourly wage for remote employees was about 60% higher than that for on-site employees. This surge could reflect changes in the dynamics of the labor market, such as increased demand for remote-capable jobs, changes in job composition, or changes in the types of occupations more suited to remote work. In 2022, the wage ratio showed a slight decline, but

it remained elevated compared to the pre-pandemic levels, suggesting a lasting impact on wage structures between remote and non-remote workers.

Why is the average real hourly wage higher for remote employees? According to the compensating wage differential theory, if working from home is considered an amenity, then wages should be lower for remote employees. (Mas and Pallais 2017) found that, on average, job applicants are willing to accept an 8 percent reduction in wages for the opportunity to work from home.

In contrast, the productivity argument suggests that wages are positively correlated with employee productivity. If working from home increases employees' labor productivity, then the wage rate should be higher for remote employees compared to on-site employees. Productivity improvements may stem from factors such as the elimination of commuting and fatigue from traffic, as well as reduced social distractions from chatting with co-workers and colleagues. (Bloom et al. 2015) conducted a randomized WFH experiment at a Chinese travel agency called Ctrip and found that working from home led to a 13% increase in performance. (Davis, Ghent, and Gregory 2021) have found that relative total factor productivity (TFP) of WFH must have increased by 48% for low-skill workers and 82% for high-skill workers to rationalize the fourfold increase in days worked from home after the COVID-19 pandemic. This could explain the increasing gap in wages for remote and non-remote employees, but it doesn't explain the existence of this gap.

To investigate the potential reasons for the observed WFH wage premium and its relationship with worker skills and occupational characteristics, we develop a structural model based on the directed search framework (Menzio and Shi 2010). This framework is particularly well-suited to our research question because it explicitly models the labor market as a collection of submarkets defined by the utility (incorporating both wage and remote work flexibility) offered to workers. By incorporating heterogeneity in both worker skill and firm remote work efficiency, the model allows us to analyze how these factors interact with firms' optimal choices of utility offers and workers' optimal search strategies. This structure enables us to explore how sorting dynamics – the matching of specific worker types with specific firm types in particular submarkets – can generate wage differentials and remote work patterns consistent with the empirical evidence, providing insights beyond reduced-form regressions.

In our study, we first empirically test the hypothesis that a significant positive difference exists between wages for remote and on-site employees. We refer to this positive gap as the work-from-home premium. Using data from the American Community Survey (2013-2023), we show that,

after controlling for education, age, race, industry, occupation, year and place of residence, the WFH premium is around 10%.

Our model analyzes how remote work affects wages using a **directed search framework**, inspired by (Menzio and Shi 2010). In this setup, the labor market isn't a single pool; instead, it's divided into many "submarkets," each defined by the specific level of utility (a combination of wage and remote work flexibility) offered to a particular type of worker. Workers, knowing their own skill level, choose which submarket (which utility offer) to search in to maximize their chances of getting a good job quickly. Firms, knowing the type of worker they want to hire, decide which utility level to promise in their job postings to attract those workers efficiently. This directed search approach allows us to explicitly model how different firms compete for different workers by offering distinct contracts (utility levels) and how this competition shapes wages and job-finding rates across the market, reflecting the idea that workers actively seek out jobs that best match their preferences and skills, rather than randomly bumping into employers.

Our study relates to two strands of literature. The COVID-19 pandemic and stay-at-home restrictions have sparked discussions on various topics, including the impact of working from home (WFH) on cities, work arrangements, and the housing market. Our paper focuses on studies examining WFH trends. In 2023, the share of full days worked from home was four times greater than in the year before the pandemic of 2019 (Barrero, Bloom, and Davis 2023). Barrero, Bloom, and Davis (2023) discuss how COVID-19 accelerated the shift to WFH and why it is likely to persist. (Davis, Ghent, and Gregory 2021) attribute the increase in WFH days to a significant rise in the productivity of WFH technology—48% for low-skill workers and 83% for high-skill workers.

Another strand of literature explores the relationship between remote work arrangements and wages. (Gariety and Shaffer 2007) empirically test the existence of a wage differential associated with working from home and find that it is associated with higher wages than traditional on-site work. (Mas and Pallais 2017) estimated employees' willingness to pay for alternative work arrangements, such as flexible scheduling, working from home, and positions with employer discretion over scheduling. They conducted a discrete choice experiment during the job application process for a national call center and found that working from home was the most valued option among employees. On average, job applicants were willing to accept a 8% reduction in wages to take advantage of the opportunity to work from home.

Our study aims to contribute to this literature by examining the potential reasons for the wage gap

between remote and non-remote employees.

The remainder of this paper is organized as follows. Section 2 establishes the empirical motivation for our study, presenting stylized facts about remote work, wages, and skills using data from the ACS, O*NET, and BLS, and highlighting the correlation between remote work/teleworkability and higher wages. Section 3 develops our theoretical framework, a directed search model incorporating heterogeneity in worker skill (h) and firm remote work efficiency (ψ), and derives the optimal remote work policy. Section 4 details the estimation and calibration strategy, including the construction of our Teleworkability Index, the econometric estimation of production function components, and the calibration of the model’s remaining parameters and distributions using external data and literature values. Section 5 presents the key qualitative results from simulating the calibrated model, focusing on its predictions for sorting, wages, remote work patterns, and market outcomes. Section 6 concludes. Technical derivations and the equilibrium computation algorithm are detailed in the Appendix.

2 Empirical Motivation: Remote Work, Wages, and Skills

Before presenting our structural model, we establish key empirical patterns using individual-level data from the American Community Survey (ACS) for 2013-2022, combined with occupation-level data from O*NET and BLS. Our sample includes civilian wage-employed individuals aged 22-70 working over 30 hours per week above the federal minimum wage. We identify remote workers based on the ACS question regarding commuting, where respondents indicate they “worked from home.”

2.1 Worker Characteristics

As shown in Table 1, remote workers differ significantly from their non-remote counterparts. Pre-pandemic (2013-2019), remote work was rare (3%), but its share increased substantially post-2020 (15%). Remote workers, on average, are similar in age but report considerably higher total income and hourly wages (both nominal and real). They are also significantly more likely to hold college (66% vs. 39%) or postgraduate (26% vs. 15%) degrees. These descriptive statistics highlight selection into remote work based on observable characteristics, particularly education and earnings potential.

Table 1: Summary statistics by work arrangement

	Non-WFH		WFH	
	Mean	Sd	Mean	Sd
Share labor force 2013 - 2019	97%	-	3%	-
Share labor force 2020 - 2023	85%	-	15%	-
Age	44.20	12.43	44.66	11.88
Total income	67,536.4	69,200.87	106,556.2	97,919.89
Hourly wage	27.95	25.81	44.20	37.59
Real hourly wage	26.31	24.10	39.31	33.10
Commuting time	26.81	23.50	-	-
Share of College	39%	-	66%	-
Share of Postgraduate	15%	-	26%	-
Observations	9025857		751654	

2.2 Stylized Fact I: Remote Work Correlates with Higher Wages, Even Controlling for Teleworkability

A simple comparison suggests remote workers earn more. However, occupations differ in their suitability for remote work. We construct a novel, data-driven **Teleworkability Index** for occupations using machine learning on high-dimensional O*NET data, validated against BLS data (see Section 4.2 for details). This index measures the feasibility of performing an occupation remotely.

Table 2 presents regressions of log real hourly wages on our Teleworkability Index and a WFH indicator. Column (1) shows a strong positive correlation between teleworkability and wages. Adding industry fixed effects (Column 2) and individual controls including detailed age-education interactions (Column 5) attenuates but does not eliminate this relationship. Crucially, when including both the Teleworkability Index and the WFH indicator (Column 6), both remain highly significant. Workers in more teleworkable occupations earn more, and *within* those occupations (controlling for teleworkability and other factors), those actually working remotely earn an additional premium (approx. 7.9%). This suggests both occupational sorting and a distinct remote work effect contribute to wage differences.

2.3 Stylized Fact II: Within Occupations, Remote Workers Earn More

To isolate the WFH premium from occupational characteristics like teleworkability or skill requirements, we regress log real hourly wages on the WFH indicator while including increasingly stringent fixed effects. Table 3 shows the results. Controlling for year, location, industry, demographics, and

Table 2: Wage regressed on Teleworkability index and remote work indicator.

	(1)	(2)	(3)	(4)	(5)	(6)
Teleworkability Index	33.58*** (0.0522)	27.68*** (0.0580)	20.89*** (0.115)	19.70*** (0.115)	15.36*** (0.0569)	14.90*** (0.0570)
WFH Indicator				6.365*** (0.0506)		3.203*** (0.0487)
FE: Year & Location	✓	✓	✓	✓	✓	✓
FE: Industry		✓		✓	✓	✓
AgeCat \times Educ					✓	✓
N	9708029	9708029	9708029	9708029	9708028	9708028
R^2	0.141	0.186	0.227	0.230	0.292	0.293

Worker level data. All regressions include demographic controls: age, race, education, others.
Standard errors in parentheses: * $p < 0.1$, ** $p < 0.05$, *** $p < 0.001$.

age-education interactions (Column 4) yields a WFH premium of approximately 13%. Adding detailed occupation fixed effects (Column 5), which absorbs time-invariant differences across occupations including average skill and teleworkability, reduces the premium to a still statistically and economically significant 8.8%. This confirms that even when comparing workers within the same detailed occupation, those working remotely tend to earn more.

Table 3: Wage regressed on remote work indicator and controls.

	(1)	(2)	(3)	(4)	(5)
WFH Indicator	12.44*** (0.0530)	7.702*** (0.0525)	5.834*** (0.0494)	5.031*** (0.0493)	3.603*** (0.0471)
FE: Year & Location	✓	✓	✓	✓	✓
FE: Industry		✓		✓	✓
FE: Occupation			✓	✓	✓
FE: Class of Worker				✓	✓
AgeCat \times Educ					✓
N	9712293	9712293	9712293	9712293	9712292
R^2	0.0711	0.153	0.289	0.307	0.364

Worker level data. All regressions include demographic controls: age, race, education, others.
Standard errors in parentheses: * $p < 0.1$, ** $p < 0.05$, *** $p < 0.001$.

2.4 Summary and Motivation for Model

The empirical evidence highlights several key facts: remote workers are positively selected on skill and earnings; occupations amenable to remote work command higher wages overall; a significant wage premium exists for remote workers even within detailed occupations; and this premium has evolved over time. While regressions control for observables, they cannot fully account for unobserved heterogeneity (e.g., firm productivity, worker preferences) or the general equilibrium effects of remote work adoption. These facts motivate our structural search and matching model, which incorporates heterogeneity in worker skill (h) and firm remote work efficiency (ψ) to explore the

mechanisms driving these wage patterns and the sorting of workers and firms in the labor market.

3 Model

3.1 Model Overview

This model follows a directed search framework in the spirit of (Menzio and Shi 2010). The model incorporates two key sources of heterogeneity: firms differ in their remote-work efficiencies, while workers vary in their skill levels. The key mechanisms in this framework are that workers value the flexibility provided by remote work arrangements, high-skilled workers are more productive and better suited for remote work, and firms treat remote and on-site work as substitutable inputs in their production processes.

Workers are characterized by their productivity h . They incur disutility from on-site work, which is partially compensated by their wage. We denote by $u(w, \alpha)$ the utility of a worker earning wage w with remote work arrangement $\alpha \in [0, 1]$, where α represents the fraction of time working remotely. We assume that utility continuously differentiable, concave and increasing in consumption $u_w(\cdot) > 0$ and in remote work $u_\alpha(\cdot) > 0$. We further assume that the marginal rate of substitution between remote work and consumption, $-u_\alpha(w, \alpha)/u_w(w, \alpha)$ is increasing in α . This condition means that as workers spend more time working remotely, they place an increasing relative value on additional remote work compared to higher wages. Intuitively, this captures increasing marginal willingness to trade off wage for flexibility: the more accustomed a worker becomes to remote work, the more they value the ability to work remotely even further. We assume that workers supply their unit of labor inelastically.

Firms are characterized by their remote-work efficiency parameter ψ , which determines how effectively they can implement remote work arrangements. The output of a firm-worker match depends on firm remote productivity, worker skill and the fraction of remote work determined by the arrangement α . We assume that $A(h)$, ($A'(h) > 0$) captures the contribution to output from worker skill h . The firm split worker's labor between in-person work and remote work. However, these components aren't perfectly substitutable, remote work is adjusted by a factor $g(\psi, h)$. The function $g(\psi, h)$ captures the efficiency of remote work, with $g_\psi(\psi, h) \geq 0$ indicating that remote work efficiency increases with firm type ψ due to better technology, management practices, or the nature of the occupation. Similarly, $g_h(\psi, h) \geq 0$ suggests that remote work efficiency increases with

worker skill level due to greater autonomy and technological ability. The production function of the firm is given by:

$$Y(\alpha \mid \psi, h) = A(h) ((1 - \alpha) + \alpha g(h, \psi)) \quad (1)$$

Profits are determined by the difference between the output produced and the wage paid to the worker. Let x denote the total utility level that a working arrangement delivers to the worker, x is derived from both the wage received and the work arrangement (i.e., the fraction of remote work α). Since workers care only about their total utility, the firm is constrained to ensure that the worker obtains at least the utility level x . This creates a trade-off for the firm: while offering higher remote work (a larger α) may decrease productivity it might also allow the firm to pay a lower wage to meet the utility guarantee. A type ψ firm with a worker of type h chooses α to maximize the following expression:

$$\Pi(\alpha \mid \psi, h, x) = \max_{\alpha \in [0,1]} \{Y(\alpha \mid \psi, h) - w(\alpha) \mid x = u(w(\alpha), \alpha)\} \quad (2)$$

Since $g(\psi, h)$ captures the adjustment of productivity to remote work in Equation 1 and is heterogeneous across firms and workers, the optimal remote work policy will also be heterogeneous.

3.2 Optimal Remote Work Choice

The firm's problem implies an optimal choice of remote work $\alpha^*(\psi, h, x)$ that satisfies the first-order condition:

Interior Solution: Let's start with the interior solution, where the firm chooses a positive fraction of remote work. This optimal choice $\alpha^*(\psi, h, x)$ must satisfy: 1. **First Order Condition:**

$$A(h) (g(\psi, h) - 1) = -\frac{u_\alpha(w(\alpha^*), \alpha^*)}{u_w(w(\alpha^*), \alpha^*)}. \quad (3)$$

2. **Promise Keeping Constraint:**

$$x = u(w(\alpha^*), \alpha^*). \quad (4)$$

The assumptions $u_w > 0$ and $u_\alpha > 0$ ensure that the utility function is strictly increasing in the wage argument and continuously differentiable, which in turn guarantees that the inversion

$w(\alpha) = u^{-1}(x, \alpha)$ is well defined and unique. Details of the derivation are on Appendix [A.1](#)

Economic Interpretation: The term $A(h)(g(\psi, h) - 1)$ represents the marginal change in production when increasing the share of remote work α . In particular, if remote work is less productive (i.e., $g(\psi, h) < 1$), then increasing α improves output; otherwise, the opposite is true. The term u_α/u_w represents the trade-off from the worker's utility perspective, indicating how much the wage can be reduced (in marginal terms) for an increase in α while still maintaining the worker's promised utility level x . The firm chooses α such that the marginal benefit from the production side exactly offsets the marginal saving in wage cost, as captured by the worker's indifference curve. Importantly, no assumption is made the magnitude of $g(\psi, h)$. Should remote work be more productive for any worker skill level ($g(\psi, h) > 1$), the firm optimally selects a corner solution with $\alpha = 1$. This arises because the marginal benefit from increasing α includes both enhanced production *and* potential wage savings ($u_\alpha/u_w > 0$). Consequently, the trade-off described earlier—balancing potential output changes against wage cost adjustments—no longer constrains the firm from maximizing remote work.

Corner Solution: Notice that Equation [3](#) describes only the interior solutions. The firm may also choose to offer either full remote work ($\alpha = 1$) or no remote work ($\alpha = 0$). The conditions for these corner solutions are as follows:

- $[\underline{\psi}(h)]$ *Threshold for Offering Some Remote Work* ($\alpha^* > 0$):

$$g(\underline{\psi}(h), h) = 1 - \frac{1}{A(h)} \left[\frac{u_\alpha(w(0), 0)}{u_w(w(0), 0)} \right] \quad (5)$$

- $[\bar{\psi}(h)]$ *Threshold for Full Remote Work* ($\alpha^* = 1$):

$$g(\bar{\psi}(h), h) = 1 - \frac{1}{A(h)} \left[\frac{u_\alpha(w(1), 1)}{u_w(w(1), 1)} \right] \quad (6)$$

Both thresholds are obtained by evaluating

$$\left. \frac{d\Pi(\alpha)}{d\alpha} \right|_{\alpha=0} \quad \text{and} \quad \left. \frac{d\Pi(\alpha)}{d\alpha} \right|_{\alpha=1} \quad (7)$$

the idea is to find for each skill level the region in the remote productivity space where even a slight deviation in remote work (from full remote or full in person) do not increase profits from the firm. Details on Appendix [A.1](#)

The optimal remote work policy is given by:

$$\alpha^*(\psi, h, x) = \begin{cases} 0 & \text{if } \psi \leq \underline{\psi}(h) \\ \alpha^*(\psi, h, x) & \text{if } \underline{\psi}(h) < \psi < \bar{\psi}(h) \\ 1 & \text{if } \bar{\psi}(h) \leq \psi \end{cases} \quad (8)$$

The interior solution $\alpha^*(\psi, h, x)$ satisfies the first-order condition in Equation 3.

Properties of the Optimal Remote Policy

1. If an increase in worker skill h raises remote productivity by more than it raises baseline productivity. In other words, then the interior solution $\alpha^*(h)$ increases with h (firms offer more remote work to more skilled workers). At the same time, the thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$, decrease with worker skill; meaning that the requirements on firm remote productivity for offering some (and full) remote work are lower.
2. If an increase in worker skill h raises remote productivity and baseline productivity equally at the margin, then the interior solution $\alpha^*(h)$ is constant with respect to h . The thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$ are locally invariant (flat) with respect to worker skill.
3. If an increase in worker skill h raises baseline productivity more than remote productivity, then the interior solution $\alpha^*(h)$ decreases with h higher worker skill implies lower remote work. The thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$ increase with worker skill; meaning that the requirements on firm remote productivity for offering some (and full) remote work are higher. Details of the derivations of this properties are on Appendix A.2

3.2.1 Parametric Example

In this example, we detail a specific parameterization that illustrates the general structure of our model. The productivity function is given by

$$A(h) = A_0 + A_1 h, \quad (9)$$

with the assumption that $A'(h) = A_1 > 0$, ensuring that productivity increases with skill level h . Here, A_0 represents the base level of productivity common to all workers, while A_1 captures the incremental productivity gain associated with an increase in worker skill h . This linear form reflects the idea that more skilled workers generate higher baseline output in a proportional manner.

The effectiveness of remote work is modeled as

$$g(\psi, h) = \psi - \psi_0 + \phi \log(h) \quad \text{with } \phi \geq 0. \quad (10)$$

ψ_0 serves as a baseline adjustment factor that reflects the overall conduciveness of the technological environment to remote work. For example, improvements in communication technologies effectively lower ψ_0 , thereby reducing the productivity losses typically associated with remote work arrangements for all possible worker-firms pairs. The term $\phi \log(h)$ captures the impact of worker skill on remote work productivity. The parameter $\phi \geq 0$ measures the sensitivity to skill—higher values of ϕ indicate that increases in worker skill more strongly enhance remote work efficiency. Importantly, the logarithmic transformation implies diminishing marginal returns to skill in this context, capturing the idea that initial improvements in h yield larger efficiency gains compared to later ones. Worker utility is specified by

$$U(w, \alpha) = a_0 + a_1 w - c_0(1 - \alpha)^\chi, \quad (11)$$

with parameters satisfying $a_1 > 0$, $c_0 > 0$, and $\chi > 1$. This formulation guarantees that utility increases with the remote work share α and that the marginal trade-off ratio is rising in α .

Given these functional forms, a higher worker skill h enhances both baseline productivity and the effectiveness of remote work. Consequently, as h increases, the model predicts a higher optimal remote work share, α^* , under sufficiently efficient remote work conditions. Calibrating parameters such as A_0 , A_1 , and ϕ enables us to match observed patterns in remote work adoption: a higher ϕ would suggest that increases in worker skill substantially lower the efficiency threshold required for firms to adopt remote work, shifting the choice of α^* upward.

The firm's profit maximization problem, subject to a fixed utility level for the worker, leads to the following optimal remote work share:

$$\alpha^*(\psi, h, x) = \begin{cases} 0 & \text{if } \psi \leq \underline{\psi}(h), \\ 1 - \left[\frac{a_1 A(h)(1 - g(\psi, h))}{c_0 \chi} \right]^{\frac{1}{\chi-1}} & \text{if } \underline{\psi}(h) < \psi < \bar{\psi}(h), \\ 1 & \text{if } \bar{\psi}(h) \leq \psi, \end{cases} \quad (12)$$

With:

$$\bar{\psi}(h) = 1 + \psi_0 - \phi \log(h), \quad \text{and} \quad \underline{\psi}(h) = 1 + \psi_0 - \phi \log(h) - \frac{c_0 \chi}{a_1(A_0 + A_1 h)}. \quad (13)$$

Regarding their monotonicity in h , if $\phi = 0$ the upper threshold remains constant, but if $\phi > 0$, $\bar{\psi}(h)$ is strictly decreasing. In contrast, with $\phi = 0$, the lower threshold $\underline{\psi}(h)$ increases with h . When $\phi > 0$, the monotonicity of $\underline{\psi}(h)$ depends on the value of A_0 . If $A_0 = 0$, $\underline{\psi}(h)$ increases for $h < \hat{h}$ and decreases for $h > \hat{h}$, where

$$\hat{h} = \frac{c_0 \chi}{\phi a_1 A_1}. \quad (14)$$

For $A_0 > 0$, if $c_0 \chi \leq 4\phi a_1 A_0$, $\underline{\psi}(h)$ is strictly decreasing; if $c_0 \chi > 4\phi a_1 A_0$, it decreases for $h \in (0, \hat{h}_1)$, increases for $h \in (\hat{h}_1, \hat{h}_2)$, and decreases again for $h > \hat{h}_2$, with the turning points determined by

$$\hat{h}_{1,2} = \frac{(c_0 \chi - 2\phi a_1 A_0) \pm \sqrt{c_0 \chi (c_0 \chi - 4\phi a_1 A_0)}}{2\phi a_1 A_1}. \quad (15)$$




Figure 1: Lower ($\underline{\psi}(h)$) and upper ($\bar{\psi}(h)$) thresholds variation with worker skill.

Figure 1 provides a graphical representation of the thresholds for a parametrization with $A_0 = 0$. The declining curves of $\underline{\psi}(h)$ and $\bar{\psi}(h)$ with higher worker skill indicate that, as employees become more skilled, firms are more likely to adopt remote work policies even when their efficiency levels are moderate. This further underscores the importance of skill in determining the mix of remote versus in-person work.




Figure 2: Optimal remote work share (α^*) for different remote efficiency and skill levels.

Figure 2, shows the optimal remote work share α^* as a function of firm remote work efficiency ψ across different worker skill levels h . The graph delineates regions where the optimal α^* takes on the values 0 (in-person work), between 0 and 1 (partial remote work), and 1 (full remote work). Notably, higher worker skill shifts the transition thresholds, implying that more skilled workers require lower firm efficiency to justify a shift towards remote work. The chosen parametrization guarantees that the interior solution $\alpha^*(h)$ increases with h .

3.3 Labor Market Search

Both firms and workers discount the future at rate β . Workers are characterized by their type h and direct their search toward submarkets distinguished by the promised utility level x . A worker of type h evaluates the different utility promises available in each submarket and chooses to search in the one that maximizes their expected value. This expected value incorporates not only the probability of being hired but also the future discounted value of the job. At the same time, firms target workers of a particular type h by posting job offers (or contracts) that promise a specific utility level x . This setup allows the market to be segmented into different submarkets. The tightness of a submarket (h, x) is defined as:

$$\theta(h, x) = \frac{v(h, x)}{u(h, x)},$$

where $v(h, x)$ denotes the number of vacancies posted by firms in the submarket and $u(h, x)$ represents the number of unemployed workers actively searching within that particular submarket. This measure of tightness directly influences the probabilities of matching: the vacancy filling rate $q(\theta(h, x))$ and the job finding rate $p(\theta(h, x))$ are both functions of θ . Matches are exogenously broken at a rate δ . Free entry of firms, incurring a cost $\kappa(x)$ per vacancy posted (which may depend on the offered utility x), ensures that market opportunities yielding positive expected profits are exploited.

Once a firm and a worker are matched, the firm delivers the promised utility x to the worker by applying the firm's optimal remote work policy $\alpha^*(\psi, h, x)$. Crucially, firms know their remote-work efficiency parameter ψ *before* deciding where to post vacancies. The distribution of firm types in the overall economy is given by $F(\psi)$, which is common knowledge.

Firm Value Functions and Probabilistic Choice:

Firms choose strategically which submarket (h, x) to target based on the potential profitability of the match. The value of an ongoing match for a firm of type ψ with a worker of type h under contract x , denoted $J(\psi, h, x)$, is determined by the current payoff plus the discounted expected continuation value:

$$J(\psi, h, x) = Y(\alpha^*(\psi, h, x)|\psi, h) - w(x, \alpha^*(\psi, h, x)) + \beta[(1 - \delta)J(\psi, h, x) + \delta V(\psi, h)], \quad (16)$$

where $V(\psi, h)$ represents the expected value for a firm of type ψ upon returning to the market after separation. This value reflects the optimal posting strategy, considering the expected returns from different submarkets.

For a given worker type h and observed market vacancy filling rates $\{q(h, x')\}_{x'}$, a firm ψ evaluates the expected profit from posting in each submarket x' :

$$\Pi_{post}(\psi, h, x') = q(h, x')J(\psi, h, x') - \kappa(x').$$

Instead of deterministically choosing only the submarket(s) yielding the absolute maximum profit, we adopt a **probabilistic choice framework** (akin to a random utility model). Firms choose submarket x with a probability $P(x|\psi, h)$ that is increasing in the potential profit $\Pi_{post}(\psi, h, x)$. A

common specification (logit choice, derived from Type I Extreme Value shocks) gives:

$$P(x|\psi, h) = \frac{\exp(\xi \cdot \Pi_{post}(\psi, h, x))}{\sum_{x''} \exp(\xi \cdot \Pi_{post}(\psi, h, x''))}, \quad (17)$$

where $\xi \geq 0$ is a parameter governing the sensitivity of choices to profit differences. As $\xi \rightarrow \infty$, the choice becomes deterministic (only the maximum profit option is chosen); as $\xi \rightarrow 0$, choices become random across profitable options. Firms only participate (i.e., $P(x|\psi, h) > 0$ for some x) if the maximum potential profit $\max_{x'} \Pi_{post}(\psi, h, x')$ is non-negative.

This probabilistic choice leads to **endogenous sorting** where firms with characteristics (ψ, h) that yield higher profits in certain submarkets x are more likely to post vacancies there.

Equilibrium and Free Entry: In equilibrium, firms' expectations about vacancy filling rates $q(h, x)$ are consistent with the rates generated by the market tightness $\theta(h, x)$ resulting from firm and worker decisions. Free entry ensures that, in active markets, expected profits are driven down. Specifically, the equilibrium vacancy filling rate $q(h, x)$ adjusts such that the *expected* profit for a firm entering market (h, x) is zero on average, considering the mix of firms likely to choose that market. This implies:

$$q(h, x) \mathbb{E}[J(\psi, h, x) \mid \text{firm } \psi \text{ chooses } (h, x)] = \kappa(x). \quad (18)$$

The conditional expectation is calculated based on the choice probabilities and the underlying firm distribution:

$$\mathbb{E}[J(\psi, h, x) \mid \text{firm } \psi \text{ chooses } (h, x)] = \frac{\sum_{\psi'} J(\psi', h, x) P(x|\psi', h) f(\psi')}{\sum_{\psi''} P(x|\psi'', h) f(\psi'')}, \quad (19)$$

where the denominator represents the total probability mass of firms choosing market (h, x) . A market (h, x) is active only if this expected value is sufficient to cover the cost $\kappa(x)$ (i.e., $\mathbb{E}[J|h, x] > \kappa(x)$). Infra-marginal firms (those with high $J(\psi, h, x)$ relative to the average) may still perceive a positive expected profit from participating.

The distribution of firms encountered by workers searching in submarket (h, x) is the *conditional distribution*, reflecting the probabilistic choices. The conditional probability density function (PDF) is:

$$f(\psi \mid h, x) = \frac{P(x|\psi, h) f(\psi)}{\sum_{\psi'} P(x|\psi', h) f(\psi')} \quad (20)$$

where the denominator is the total mass choosing (h, x) . The conditional cumulative distribution function (CDF) $F(\psi | h, x)$ is obtained by integrating (or summing) this density:

$$F(\psi | h, x) = \int_{-\infty}^{\psi} f(\psi' | h, x) d\psi' = \frac{\sum_{\psi'' \leq \psi} P(x|\psi'', h) f(\psi'')}{\sum_{\psi'} P(x|\psi', h) f(\psi')} \quad (21)$$

Worker Value Functions: For workers, the value functions capture the trade-off between being unemployed and employed. A worker of type h chooses the submarket x that maximizes their expected lifetime utility. The value of unemployment $U(h)$ is:

$$U(h) = b + \beta \max_x \left\{ p(\theta(h, x)) \mathbb{E}_{\psi|h, x} [W(\psi, h, x)] + (1 - p(\theta(h, x))) U(h) \right\}, \quad (22)$$

where $p(\theta(h, x))$ is the job finding rate, and the expectation $\mathbb{E}_{\psi|h, x}[\cdot]$ uses the conditional distribution $F(\psi|h, x)$ from Equation 21:

$$\mathbb{E}_{\psi|h, x} [W(\psi, h, x)] = \int W(\psi', h, x) dF(\psi'|h, x) = \sum_{\psi'} W(\psi', h, x) f(\psi'|h, x).$$

This reflects that workers anticipate the specific mix of firms attracted to each submarket (h, x) .

Once employed with a firm of type ψ under contract x , the worker's value $W(\psi, h, x)$ evolves according to:

$$W(\psi, h, x) = x + \beta \left[(1 - \delta) W(\psi, h, x) + \delta U(h) \right]. \quad (23)$$

The equilibrium vacancy filling rates $q^*(h, x)$, market tightnesses $\theta^*(h, x)$, firm choice probabilities $P^*(x|\psi, h)$, and the resulting conditional distribution of firms $f^*(\psi|h, x)$ are determined simultaneously. These objects satisfy the firm's probabilistic choice rule Equation 17, the free entry condition Equation 18 based on the conditional expectation Equation 19, and consistency with the underlying matching technology $q(\theta)$. We solve for this equilibrium numerically using an iterative algorithm detailed in Appendix A.3.

4 Calibration

This section details the procedure used to discipline the model and obtain plausible values for key structural parameters governing firm production, worker skills, and firm efficiency distributions.

Our approach combines econometric estimation using detailed industry-occupation data with calibration techniques to map estimated coefficients to the model’s parameters. We leverage data from the Bureau of Labor Statistics (OES, Productivity), the American Community Survey (ACS), the Census Bureau’s Business Dynamics Statistics (BDS), and the O*NET database.

4.1 Data Sources and Variable Construction

We construct a panel dataset primarily at the 3-digit NAICS industry level, spanning the period from 2013 to 2023 . Key variables include:

- **Output per Hour:** Derived from BLS sectoral productivity data, adjusted for price deflation. This serves as our measure of Y/L in the data.
- **Industry Skill Index:** Constructed by weighting occupation-level skill indices derived from O*NET data (combining normalized importance and level ratings for relevant skills and abilities) by their employment shares within each industry, using BLS OES data. This proxies for the average h within an industry.
- **Industry Teleworkability Index:** Self-constructed index based on the Teleworkability Index at the occupation level. The procedure is described in the next section, weighted by employment shares within industries.
- **Observe Remote Work Rates:** Measured using data from the ACS, representing the observed share of remote work in each industry-year.

4.2 Estimating the Occupation Teleworkability Index

A key input into our structural model is the heterogeneity in firms’ remote work efficiency (ψ). We require an empirical counterpart to discipline the distribution of this parameter. Existing indices of occupation-level remote work feasibility, such as the binary classification by (Dingel and Neiman 2020) or the index by (Mongey, Pilossoph, and Weinberg 2020), provide valuable benchmarks. These approaches typically involve selecting key occupational characteristics based on theoretical considerations or researcher assessment to determine teleworkability. To leverage a richer, data-driven approach using a high-dimensional feature set, we construct a novel **Teleworkability Index** using machine learning techniques on detailed occupational data. This index aims to capture both whether an occupation *can* be performed remotely (extensive margin) and the *degree* to which it is feasible (intensive margin).

Data: Our primary feature set consists of detailed occupation-level characteristics from the O*NET database, encompassing a wide range of variables related to skills, abilities, work activities, and work context. For supervised learning, we utilize data from the Occupational Requirements Survey, which reports the share of workers in each occupation for whom telework is available. In 2024, telework was available to 14.9% of civilian workers. We normalize this variable to a $[0, 1]$ scale. As ORS coverage is incomplete, this provides labels for a subset of occupations.

Methodology: We employ a two-stage modeling approach using Random Forest models, leveraging their ability to handle high-dimensional data and capture non-linear relationships. Both models' hyperparameters are tuned using cross-validation techniques on the training data.

1. **Stage 1: Classification (Extensive Margin):** We first train a *Random Forest Classifier* to distinguish occupations where telework is impossible (ORS label = 0) from those where it is at least partially possible (ORS label > 0). The model is trained on the labeled subset of occupations using the binary indicator derived from the ORS label. This stage identifies occupations fundamentally unsuited for remote work, which is particularly important because of the large amounts of zeros in the training data.
2. **Stage 2: Regression (Intensive Margin):** For occupations identified as having *non-zero* teleworkability by the ORS labels, we train a *Random Forest Regressor*. To ensure predictions remain within the $[0, 1]$ bound and to better handle the distribution of the target variable, we apply a logit transformation ($\log(y/(1-y))$) to the non-zero ORS labels before training. The regressor learns the relationship between ONET features and the *degree* of teleworkability for jobs where remote work is feasible.

Validation and Prediction: The performance of the tuned two-stage model is evaluated using bootstrap validation on held-out test data, assessing metrics such as Accuracy and F1-score for the classifier, and Mean Squared Error (MSE) and Correlation for the regressor. The classifier achieved an accuracy of 90.7% along with an F1 score of 91.9%, while the regression model recorded a mean squared error (MSE) of 0.1 and a correlation of 0.71; Figure 3 shows the distribution of predicted teleworkability indices for both labeled and unlabeled data. The final trained model is then used to predict the teleworkability index for *all* occupations in the O*NET database, including those without ORS labels. For a given occupation:

- The classifier predicts the probability of it being a zero-teleworkability job. If this probability exceeds a threshold, the index is set to 0.

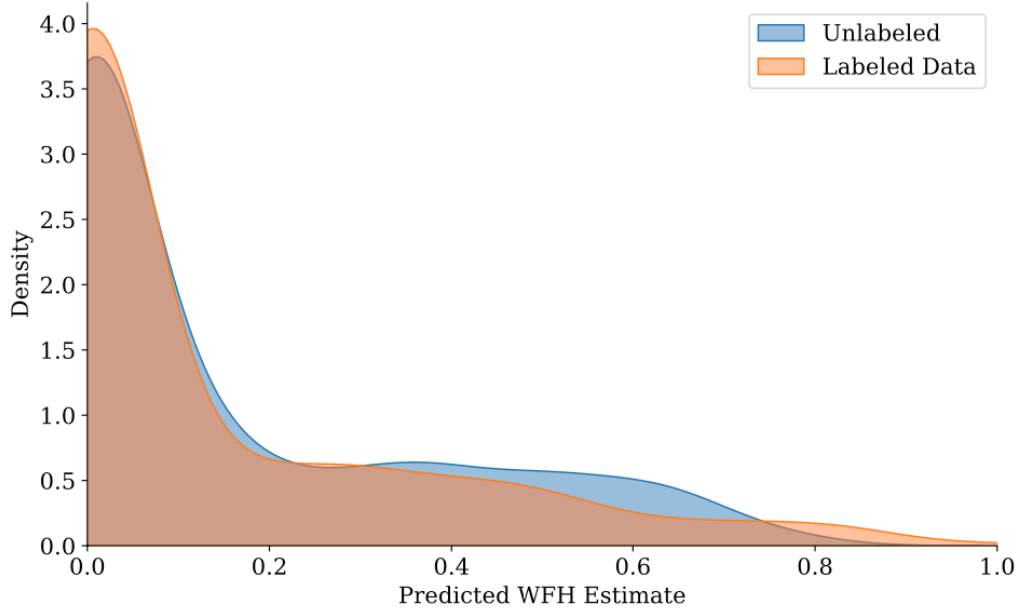


Figure 3

- Otherwise, the regressor predicts the logit-transformed level of teleworkability, which is then converted back to the $[0, 1]$ scale using the inverse logit function.

This resulting continuous index represents our measure of occupation-level teleworkability.

Linking to the Model: This estimated occupation-level Teleworkability Index forms the basis for constructing our industry-level TELE variable (by weighting occupational indices by employment shares within industries). Furthermore, the relationship estimated that we estimate allows us to use the distribution of this index across industries (weighted by firm creation) to empirically inform the aggregate distribution of firm remote work efficiency, $F(\psi)$, that we supply externally to the structural model.

4.3 Estimation of Production Function Components

To inform the parameters of the model's production function, $Y = A(h)((1 - \alpha) + \alpha g(\psi, h))$, we estimate a reduced-form relationship between productivity, skills, remote work, and teleworkability using our panel data. We assume that the industry-level data reflects the aggregate outcomes of the firm-level production specified in the model. Specifically, we relate the firm's remote work efficiency $g(\psi, h)$ to the observed industry teleworkability (TELE) and potentially worker skill (SKILL), and the baseline productivity $A(h)$ to worker skill (SKILL).

Our preferred specification is:

$$\begin{aligned}
\log(\text{OUTPUT}_{it}) = & \beta_0 \cdot \text{SKILL}_{it} \\
& + \beta_1 \cdot (\text{ALPHA}_{it} \times \text{SKILL}_{it}) \\
& + \beta_2 \cdot (\text{TELE}_{it} \times \text{ALPHA}_{it} \times \text{SKILL}_{it}) \\
& + \beta_3 \cdot (\text{SKILL}_{it} \times \log(\text{SKILL}_{it})) \\
& + \mu_i + \lambda_t + \epsilon_{it}
\end{aligned} \tag{24}$$

where i denotes industry and t denotes year, μ_i are sector fixed effects, and λ_t are year fixed effects.

The regression results are summarized in Table 4.

	OUTPUT				
	(1)	(2)	(3)	(4)	(5)
(Intercept)	16.739** (6.424)				
SKILL	25.708*** (2.353)	23.910*** (2.266)	26.279*** (2.495)	37.024*** (2.922)	39.948*** (3.158)
ALPHA \times SKILL	-37.564* (15.298)	-34.714* (15.366)	-53.873* (26.421)	-29.745 (15.311)	-84.807** (26.700)
SKILL \times log(SKILL)	64.152*** (18.611)	15.974*** (2.148)	63.632*** (18.837)	149.791*** (22.199)	157.405*** (22.526)
TELE \times ALPHA \times SKILL	315.199*** (70.115)	286.139*** (69.705)	356.831*** (89.545)	176.199* (78.090)	318.733** (96.738)
YEAR Fixed Effects			Yes		Yes
SECTOR Fixed Effects				Yes	Yes
N	418	418	418	418	418
R^2	0.263	0.510	0.265	0.378	0.388
Within- R^2			0.264	0.292	0.303

Table 4: Regression results for industry-level output estimation

Mapping Estimates to Structural Parameters

We map the estimated coefficients from (24) (specifically, using the results from model (1)) to the

Parameter	Mapping	Value
A	β_0	25.71
ψ_0	$(\beta_1 + 1)/A$	-1.42
ψ_1	β_2/A	12.26
ϕ	β_3/A	2.50
C	(Intercept)	16.74

Table 5: Model Parameters, Mappings to Regression Coefficients, and Values

parameters of our assumed functional forms for $A(h)$ and $g(\psi, h)$. with the parametric assumption

$$Y = A(h)(1 + \alpha(g(\psi, h) - 1)) \quad \text{with} \quad g(\psi, h) \approx \psi_0 + \psi + \phi \log h \quad \text{and} \quad A(h) = A_1 h$$

We derive the structural parameters $A_1, \phi, \psi_0, \psi_1$. The estimated values used in the model simulation are presented in Table 5.

Estimating Distributions




Figure 4

- **Firm Type Distribution $F(\psi)$:** We use the relationship derived from the regression, $\psi \approx \psi_1 \times \text{TELE}_i$, to obtain an estimate of the average firm type ψ_i for each industry i . Using industry weights based on job creation from new firms (from BDS), we estimate the aggregate distribution of firm types $f(\psi)$ using Kernel Density Estimation (KDE). The resulting distribution Figure 4 captures the heterogeneity in remote work efficiency across the economy. We then discretize this distribution onto the model's ψ_{grid} .
- **Worker Skill Distribution $F(h)$:** We construct the aggregate distribution of worker skills $f(h)$ similarly. We use the occupation-level skill indices derived from O*NET and weight them by their aggregate employment. We then fit a KDE to these weighted skill indices to obtain the distribution $f(h)$ Figure 5 and discretize it onto the model's h_{grid} .

Calibration of Remaining Parameters Several model parameters are not directly identified by the estimation procedure above and are calibrated using standard values from the literature or set to match specific aggregate targets.




Figure 5

- **Discount Factor (β):** Set to 0.9615 corresponding to an annual discount rate of 4%.
- **Exogenous Separation Rate (δ):** Set to 0.07.
- **Matching Function Exponent (γ):** Set to 0.6 following (Menzio and Shi 2010)
- **Unemployment Benefit (b):** Set the unemployment benefit as 0.6 of the potential wage the lowest skill worker

5 Results

This section presents the key qualitative predictions generated by simulating our directed search model, using the parameters and distributions estimated and calibrated as detailed in Section 4. We focus on how the interplay between worker skill (h), firm remote work efficiency (ψ), utility-dependent posting costs ($\kappa(x)$), and probabilistic firm choices shapes labor market outcomes.

Worker Search Behavior: The model predicts workers' optimal search strategies across utility submarkets. Figure Figure 6 plots the optimal promised utility level $x^*(h)$ targeted by workers of different skill levels h . Consistent with intuition, higher-skilled workers, who generate more value in matches, direct their search towards submarkets offering higher lifetime utility. This reflects their ability to command better compensation packages, which include both wages and remote work flexibility.

This confirms that the model generates sorting where more efficient firms compete for more skilled workers by offering more attractive utility packages.




Figure 6: Higher skill workers target higher utility submarkets.

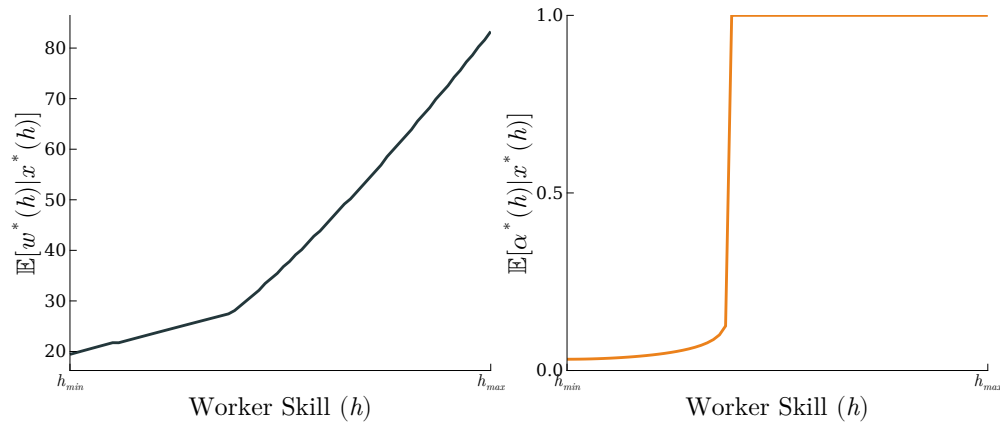


Figure 7: The model generates a positive correlation of skill with both expected wage and expected remote work share.

Expected Wages and Remote Work: The sorting patterns translate into equilibrium wage and remote work outcomes. Figure Figure 7 plots the expected wage $\mathbb{E}[w|h, x^*(h)]$ and expected remote work share $\mathbb{E}[\alpha^*|h, x^*(h)]$ for workers across the skill distribution, calculated using the conditional distribution of firms $f^*(\psi|h, x^*(h))$ in their optimally chosen submarket $x^*(h)$. As expected, higher-skilled workers command significantly higher wages. Furthermore, the expected share of remote work α^* also increases with worker skill h . This arises because higher- h workers sort into markets with higher- ψ firms, and these firms, being more efficient at remote work, optimally offer a greater share of remote time (α^*) as part of the utility package $x^*(h)$.




Figure 8

Market Outcomes: Job Finding Rates: Figure Figure 8 illustrates the equilibrium job-finding probability $p(\theta^*(h, x))$ faced by workers in different submarkets. While multiple markets might be theoretically available, workers concentrate their search in the market $x^*(h)$ offering the highest expected value. The plot shows $p(\theta^*(h, x))$ across the range of potential utility offers x for representative low, medium, and high-skilled workers. We observe that for each skill level, the job-finding probability decreases as the offered utility x increases, reflecting the higher cost and potentially lower equilibrium q^* in higher-utility markets. However, higher-skilled workers can access markets with higher utility levels before the job-finding probability drops to zero, consistent with their higher market value.

Summary of Results: The simulated model, parameterized using our empirical strategy, suc-

cessfully generates key qualitative features consistent with labor market observations. It produces positive assortative matching between worker skill and firm efficiency, mediated through directed search over utility offers. This sorting leads to higher-skilled workers targeting and receiving higher utility offers, which translates into higher expected wages and a greater expected share of remote work. The model highlights how firm heterogeneity in remote work efficiency and utility-dependent posting costs can shape equilibrium wages and work arrangements.

6 Conclusion

This paper investigates the persistent wage premium observed for remote workers in the United States. We begin by documenting key empirical facts using ACS, O*NET, and BLS data: remote workers are positively selected on earnings and education, occupations suitable for remote work command higher wages, a significant WFH premium exists even within detailed occupations, and this premium has evolved over time, increasing notably since the COVID-19 pandemic. These findings motivate our central question: what drives the wage gap between remote and non-remote employees?

To explore potential mechanisms beyond observable characteristics, we developed a directed search model featuring heterogeneity in worker skill and firm remote work efficiency. In our framework, firms strategically post utility offers (bundles of wages and remote work to attract specific worker types, incurring utility-dependent posting costs.

Workers direct their search towards utility offers that maximize their expected value. Firms make probabilistic choices over which submarket to enter, leading to endogenous sorting.

We estimate key production function parameters and distributions using industry-level data and calibrate the remaining parameters. Simulating the calibrated model reveals that it successfully generates positive assortative matching: higher-skilled workers target and match with higher-efficiency firms offering greater utility. Crucially, these higher utility packages optimally bundle both higher wages and a greater share of remote work for high-skill workers matched with high-efficiency firms.

Our results provide a structural explanation for the WFH premium, attributing it significantly to sorting dynamics and heterogeneity in firms' ability to effectively implement remote work. The model demonstrates how competition for high-skill workers leads efficient firms to offer attractive packages where remote work acts partially as an amenity but is also bundled with higher wages

reflecting the value generated in these high-productivity matches. The framework highlights that understanding the WFH premium requires considering not just worker preferences or average productivity effects, but also the equilibrium sorting outcomes driven by firm-level heterogeneity in remote work capabilities and the costs associated with different contract offers. Future work could further explore the role of worker preferences and dynamic considerations like on-the-job search.

References

- Barrero, José María, Nicholas Bloom, and Steven J Davis. 2023. “The Evolution of Work from Home.” *Journal of Economic Perspectives* 37 (4): 23–49.
- Bloom, Nicholas, James Liang, John Roberts, and Zhichun Jenny Ying. 2015. “Does Working from Home Work? Evidence from a Chinese Experiment.” *The Quarterly Journal of Economics* 130 (1): 165–218.
- Davis, Morris A, Andra C Ghent, and Jesse M Gregory. 2021. “The Work-from-Home Technology Boon and Its Consequences.” National Bureau of Economic Research.
- Dingel, Jonathan I., and Brent Neiman. 2020. “How Many Jobs Can Be Done at Home?” *Journal of Public Economics* 189 (September): 104235. <https://doi.org/10.1016/j.jpubeco.2020.104235>.
- Gariety, Bonnie Sue, and Sherrill Shaffer. 2007. “Wage Differentials Associated with Working at Home.” *Monthly Lab. Rev.* 130: 61.
- Holgersen, Henning, Zhiyang Jia, and Simen Svenkerud. 2021. “Who and How Many Can Work from Home? Evidence from Task Descriptions.” *Journal for Labour Market Research* 55 (1): 1–13.
- Mas, Alexandre, and Amanda Pallais. 2017. “Valuing Alternative Work Arrangements.” *American Economic Review* 107 (12): 3722–59.
- Menzio, Guido, and Shouyong Shi. 2010. “Block Recursive Equilibria for Stochastic Models of Search on the Job.” *Journal of Economic Theory* 145 (4): 1453–94. <https://doi.org/10.1016/j.jet.2009.10.016>.
- Mongey, Simon, Laura Pilossoph, and Alex Weinberg. 2020. “Which Workers Bear the Burden of Social Distancing?” w27085. Cambridge, MA: National Bureau of Economic Research. <https://doi.org/10.3386/w27085>.

A Appendix

A.1 Optimal Remote Policy

Problem Statement: We consider the following optimization problem:

$$\max_{\alpha \in [0,1]} \left\{ Y(\alpha \mid \psi, h) - w(\alpha) \mid x = u(w(\alpha), \alpha) \right\}$$

Invertibility Check: The key condition needed here is that the mapping $w \mapsto u(w, \alpha)$ is *strictly increasing*. Since we assume $u_w(w, \alpha) > 0$ for every (w, α) in the domain of $u(\cdot)$, this means that it can be uniquely inverted in w . Define $\omega(x, \alpha)$ as the inverse function of $u(w, \alpha)$:

$$w(\alpha) = \omega(x, \alpha) \quad \implies \quad x = u(w(\alpha), \alpha) \quad \text{for } w \in [0, \infty).$$

Continuous differentiability of u ensures that $\omega(x, \alpha)$ is well-behaved.

Substitute $w(\alpha)$ with $\omega(x, \alpha)$ in the objective function:

$$\max_{\alpha \in [0,1]} \Pi(\alpha) \quad \text{where} \quad \Pi(\alpha) = Y(\alpha \mid \psi, h) - \omega(x, \alpha)$$

Applying this substitution and considering any potential $\alpha \in \mathbb{R}$, the firm's problem becomes an unconstrained maximization in α :

$$\max_{\alpha \in \mathbb{R}} \Pi(\alpha) \quad \text{where} \quad \Pi(\alpha) = Y(\alpha \mid \psi, h) - \omega(x, \alpha)$$

The first order condition (FOC) for this maximization problem is given by:

$$\frac{d\Pi(\alpha)}{d\alpha} = A(h) (g(\psi, h) - 1) - \frac{\partial \omega(x, \alpha)}{\partial \alpha} = 0$$

Recall:

$$x = u(\omega(x, \alpha), \alpha).$$

Differentiate both sides with respect to α :

$$0 = \frac{d}{d\alpha} u(\omega(x, \alpha), \alpha)$$

Applying the chain rule:

$$0 = u_w(\omega(x, \alpha), \alpha) \frac{\partial \omega(x, \alpha)}{\partial \alpha} + u_\alpha(\omega(x, \alpha), \alpha)$$

Solving for $\partial \omega(x, \alpha) / \partial \alpha$, we obtain:

$$\frac{\partial \omega(x, \alpha)}{\partial \alpha} = - \frac{u_\alpha(\omega(x, \alpha), \alpha)}{u_w(\omega(x, \alpha), \alpha)}$$

Substituting back into the derivative of $\Pi(\alpha)$:

$$\frac{d\Pi(\alpha)}{d\alpha} = A(h) (g(\psi, h) - 1) + \frac{u_\alpha(\omega(x, \alpha), \alpha)}{u_w(\omega(x, \alpha), \alpha)}$$

For an interior optimum α^* , the first order condition is given by:

$$A(h) (g(\psi, h) - 1) + \frac{u_\alpha(\omega(x, \alpha^*), \alpha^*)}{u_w(\omega(x, \alpha^*), \alpha^*)} = 0$$

Since by definition $\omega(x, \alpha^*) = w(\alpha^*)$ (because the substitution recovers the original wage from the constraint $x = u(w(\alpha^*), \alpha^*)$), we can equivalently write:

$$A(h) (g(\psi, h) - 1) = - \frac{u_\alpha(w(\alpha^*), \alpha^*)}{u_w(w(\alpha^*), \alpha^*)}$$

In addition, the original promise-keeping constraint still holds:

$$x = u(w(\alpha^*), \alpha^*).$$

Threshold Conditions: In order to fully characterize the solution, we now derive threshold conditions for two boundary cases: - When the firm decides to offer some remote work (i.e., an optimal $\alpha^* > 0$). - When the firm decides to go all remote (i.e., $\alpha^* = 1$). *Threshold for Offering Some Remote Work ($\alpha^* > 0$):* For the firm to choose a strictly positive level of remote work, a marginal increase in α from zero must yield an increase in profit. In other words, the derivative of $\Pi(\alpha)$ evaluated at $\alpha = 0$ must be **positive**:

$$\left. \frac{d\Pi(\alpha)}{d\alpha} \right|_{\alpha=0} = A(h)(g(\psi, h) - 1) + \frac{u_\alpha(w(0), 0)}{u_w(w(0), 0)} > 0$$

Rearranging the inequality, the threshold condition is:

$$A(h)(g(\psi, h) - 1) > -\frac{u_\alpha(w(0), 0)}{u_w(w(0), 0)}$$

If this condition holds, then even a slight increase in α from zero improves profits and the firm will offer some remote work, i.e. $\alpha^* > 0$. We can solve for the threshold $\underline{\psi}(h)$ such that:

$$g(\underline{\psi}(h), h) = 1 - \frac{1}{A(h)} \left[\frac{u_\alpha(w(0), 0)}{u_w(w(0), 0)} \right]$$

Threshold for Full Remote Work ($\alpha^ = 1$):* Similarly, for the firm to choose full remote work, the derivative of $\Pi(\alpha)$ at the upper boundary $\alpha = 1$ must be non-negative. That is,

$$\left. \frac{d\Pi(\alpha)}{d\alpha} \right|_{\alpha=1} = A(h)(g(\psi, h) - 1) + \frac{u_\alpha(w(1), 1)}{u_w(w(1), 1)} \geq 0$$

When this condition holds, increasing α further (beyond values arbitrarily close to 1) does not raise profits, so the firm optimally opts for a full remote work arrangement, i.e. $\alpha^* = 1$. Similarly, we can solve for the threshold $\bar{\psi}(h)$:

$$g(\bar{\psi}(h), h) = 1 - \frac{1}{A(h)} \left[\frac{u_\alpha(w(1), 1)}{u_w(w(1), 1)} \right]$$

A.2 Proof of Properties of the Optimal Remote Policy

Proposition A.1 (Properties of the Optimal Remote Work Policy). *Consider the optimal remote work policy*

$$\alpha^*(\psi, h, x) = \begin{cases} 0 & \text{if } \psi \leq \underline{\psi}(h), \\ \alpha^*(\psi, h, x) & \text{if } \underline{\psi}(h) < \psi < \bar{\psi}(h), \\ 1 & \text{if } \bar{\psi}(h) \leq \psi, \end{cases}$$

with the interior solution $\alpha^(\psi, h, x)$ satisfying the first-order condition*

$$A(h)(g(\psi, h) - 1) = -\frac{u_\alpha(w(\alpha^*), \alpha^*)}{u_w(w(\alpha^*), \alpha^*)}.$$

and the thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$ defined implicitly by:

$$A(h)\left(g(\underline{\psi}(h), h) - 1\right) = C_0, \quad \text{and} \quad A(h)\left(g(\bar{\psi}(h), h) - 1\right) = C_1,$$

Then the following results hold: 1. If worker skill enhances relative remote productivity:

$$\frac{\partial}{\partial h} \left[A(h) g(\psi, h) \right] > \frac{\partial A(h)}{\partial h}, \quad (25)$$

then the interior solution $\alpha^*(h)$ is strictly increasing in h . And the thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$ are strictly decreasing in h . 2. If worker skill affects remote and baseline productivity equally at the margin:

$$\frac{\partial}{\partial h} \left[A(h) g(\psi, h) \right] = \frac{\partial A(h)}{\partial h}, \quad (26)$$

then the interior solution $\alpha^*(h)$ is constant with respect to h . And the thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$ are locally invariant with respect to h . 3. If worker skill affects remote and baseline productivity equally at the margin:

$$\frac{\partial}{\partial h} \left[A(h) g(\psi, h) \right] < \frac{\partial A(h)}{\partial h}, \quad (27)$$

then the interior solution $\alpha^*(h)$ is strictly decreasing in h . And the thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$ are strictly increasing in h .

Proof. Monotonicity of the Interior Solution: We are concerned with the interior solution $\alpha^*(h)$, which satisfies the first order condition (FOC):

$$F(\alpha^*(h), \psi, h, x) \equiv A(h)(g(\psi, h) - 1) + \frac{u_\alpha(w(\alpha^*(h)), \alpha^*(h))}{u_w(w(\alpha^*(h)), \alpha^*(h))} = 0$$

Our objective is to determine the sign of $\partial\alpha^*(h)/\partial h$. Differentiating the identity $F(\alpha^*(h), \psi, h, x) = 0$ with respect to h using the chain rule yields:

$$F_h + F_\alpha \frac{\partial\alpha^*(h)}{\partial h} = 0 \quad \implies \quad \frac{\partial\alpha^*(h)}{\partial h} = -\frac{F_h}{F_\alpha}$$

Notice that:

$$-F_\alpha = -\frac{\partial}{\partial \alpha} \left[\frac{u_\alpha(w(\alpha), \alpha)}{u_w(w(\alpha), \alpha)} \right] > 0 \quad (\text{by assumption})$$

The sign of $\partial\alpha^*(h)/\partial h$ therefore depends directly on the sign of F_h . Taking the partial derivative

of F with respect to h , we obtain:

$$F_h = \frac{\partial}{\partial h} [A(h)(g(\psi, h) - 1)] + \frac{\partial}{\partial h} \left[\frac{u_\alpha(w(\alpha), \alpha)}{u_w(w(\alpha), \alpha)} \right]$$

We assume that skill h enters the firm's problem only through the production side ($A(h)$ and $g(\psi, h)$) and not through the worker's utility function or the wage determination mechanism $\omega(x, \alpha)$ directly for a given x and α . Thus, the second term involving the marginal rate of substitution is zero with respect to h : $\frac{\partial}{\partial h} \left[\frac{u_\alpha}{u_w} \right] = 0$. This simplifies F_h to:

$$F_h = \frac{\partial}{\partial h} [A(h) g(\psi, h)] - \frac{\partial A(h)}{\partial h} \quad (28)$$

From this expression is clear that the conditions in the proposition directly determine the sign of F_h thus determining the monotonicity of $\alpha^*(h)$.

Monotonicity of the Thresholds: We begin with the implicit definition of $\psi(h)$ (which may denote $\underline{\psi}(h)$ or $\bar{\psi}(h)$):

$$A(h)(g(\psi(h), h) - 1) = C,$$

where C is a constant (either C_0 or C_1). Differentiating both sides with respect to h , we have

$$\frac{d}{dh} [A(h)(g(\psi(h), h) - 1)] = A'(h)(g(\psi(h), h) - 1) + A(h)(g_\psi(\psi(h), h) \cdot \psi'(h) + g_h(\psi(h), h)) = 0.$$

Solving for $\psi'(h)$:

$$\psi'(h) = -\frac{A'(h)(g(\psi(h), h) - 1) + A(h)g_h(\psi(h), h)}{A(h)g_\psi(\psi(h), h)} = -\frac{\frac{\partial}{\partial h} [A(h) g(\psi(h), h)] - \frac{\partial A(h)}{\partial h}}{A(h)g_\psi(\psi(h), h)},$$

Because $A(h) > 0$ and $g_\psi(\psi, h) > 0$, the denominator is positive. Hence, the sign of $\psi'(h)$ is determined by the numerator. From here is clear that the conditions in the proposition directly determine the monotonicity of the thresholds $\underline{\psi}(h)$ and $\bar{\psi}(h)$. \square

A.3 Equilibrium Computation Algorithm

This appendix details the iterative algorithm used to compute the equilibrium vacancy filling rates $q^*(h, x)$, the associated market tightnesses $\theta^*(h, x)$, the firm choice probabilities $P^*(x|\psi, h)$, and the

conditional firm distributions $f^*(\psi|h, x)$ and $F^*(\psi|h, x)$ described in Section 3.3. The algorithm searches for a fixed point in the vacancy filling rates $q(h, x)$ by iterating on firm profit calculations, probabilistic choices, and the free entry condition.

- **Inputs:** Grids $\mathcal{H}, \Psi_{grid}, \mathcal{X}$; Match values $J(\psi, h, x)$; Density $f(\psi)$; Cost $\kappa(x)$; Matching functions $q(\theta), q^{-1}(q)$; Sensitivity ξ ; Tolerance ϵ ; Max iterations N_{\max} .
- **Outputs:** Equilibrium rates $q^*(h, x), \theta^*(h, x)$; Choice probabilities $P^*(x|\psi, h)$; Conditional distributions $f^*(\psi|h, x), F^*(\psi|h, x)$.

Algorithm 1 Equilibrium Computation (Probabilistic Choice)

```
1: Initialize  $q(h, x) \leftarrow q_{init}$ ,  $n \leftarrow 0$ ,  $converged \leftarrow \text{False}$ .
2: while not converged and  $n < N_{\max}$  do
3:    $n \leftarrow n + 1$ .
4:    $q_{old} \leftarrow q$ .
5:   Initialize  $P_{choice}(\psi, h, x) \leftarrow 0.0$ ,  $q_{new}(h, x) \leftarrow 0.0$ .
   // Calculate choice probabilities based on expected profit  $\Pi_{post} = q_{old}J - \kappa(x)$ 
6:   for all firm types  $\psi \in \Psi_{grid}$  do
7:     for all worker types  $h \in \mathcal{H}$  do
8:       Find  $\Pi_{max}(\psi, h) \leftarrow \max_{x'} \{\Pi_{post}(\psi, h, x')\}$ .
9:       if  $\Pi_{max}(\psi, h) \geq -\epsilon$  then
10:        Calculate  $E_{x'} \leftarrow \exp(\xi \cdot (\Pi_{post}(\psi, h, x') - \Pi_{max}(\psi, h)))$  for all  $x'$ .
11:        Calculate sum  $S \leftarrow \sum_{x''} E_{x''}$ .
12:        if  $S > \epsilon$  then  $P_{choice}(\psi, h, x') \leftarrow E_{x'}/S$  for all  $x'$ .
   // Update target  $q_{new}$  based on free entry  $q_{new} = \kappa(x)/\mathbb{E}[J|h, x]$ 
13:   for all worker types  $h \in \mathcal{H}$  do
14:     for all utility levels  $x \in \mathcal{X}$  do
15:       Calculate total mass  $M \leftarrow \sum_{\psi'} P_{choice}(\psi', h, x)f(\psi')$ .
16:       if  $M > \epsilon$  then
17:        Calculate  $\mathbb{E}[J|h, x] \leftarrow (\sum_{\psi'} J(\psi', h, x)P_{choice}(\psi', h, x)f(\psi'))/M$ .
18:        if  $\mathbb{E}[J|h, x] > \kappa(x) + \epsilon$  then
19:           $q_{target} \leftarrow \kappa(x)/\mathbb{E}[J|h, x]$ .
20:           $q_{new}(h, x) \leftarrow \max(0, \min(1, q_{target}))$ .
21:   Check convergence and update:  $q \leftarrow \text{UpdateRule}(q_{old}, q_{new})$ .

22: Post-Processing:
23:  $q^* \leftarrow q$ ;  $P^* \leftarrow P_{choice}$ . and  $\theta^*(h, x) \leftarrow q^{-1}(q^*(h, x))$ .
24: Compute  $f^*(\psi|h, x)$  and  $F^*(\psi|h, x)$  using  $P^*$  and  $f(\psi)$ .

25: return  $q^*, \theta^*, P^*, f^*, F^*$ .
```
